

Tighter Approximation Schemes for Chain-Graph Logical Credal Networks

IBM Research

Abstract

A chain-graph Logical Credal Network (LCN) admits a factorization that captures its Local Markov Condition (LMC) independence assumptions. The current `cn/` pipeline exploits this by compiling the LCN into a *credal network*: a directed graph whose local credal sets $K(\text{child} \mid \text{parents})$ are interval-valued and computed per family, then performing posterior inference (Credal Variable Elimination) over the *strong extension* of that credal network. The strong extension is generally a *proper superset* of the LCN’s distribution set, so the computed posterior marginal bounds are looser than the exact LCN bounds. This note dissects the two distinct sources of this looseness on a running example (the `alarm` network), proposes five approximation schemes that produce tighter outer approximations—hence tighter posterior bounds—illustrates each on the example, explains how the existing inference algorithms (Credal VE, Interval BP, ApproxLP, Credal CTE, Credal IJGP) all inherit the tightening, and gives a validation strategy based on certified local solves and the existing LMC factorization verifier.

1 Background and notation

Let an LCN over propositional atoms $\mathbf{V} = \{V_1, \dots, V_n\}$ define a set of joint distributions

$$\mathcal{D}_{\text{LCN}} = \{ P \in \Delta(2^{\mathbf{V}}) : P \text{ satisfies every sentence } s \},$$

where each sentence s is a Type-1 constraint $\ell \leq P(\varphi) \leq u$ or a Type-2 constraint $\ell \leq P(\varphi \mid \psi) \leq u$. When the LCN’s structure graph is a chain graph, the joint factorizes over the families (child + parents) of the simplified structure graph,

$$P(\mathbf{V}) = \prod_v P(\text{ch}_v \mid \text{pa}_v), \tag{1}$$

and this factorization reproduces exactly the LMC independencies (cf. `docs/ISIPTA2025_LCNs.pdf` Sec. 3, `docs/LCN_IJAR_Revised.pdf`).

The `cn/` pipeline realizes (1) in three stages, each of which is implemented in a distinct module:

1. **Symbolic factorization** (`cn/factorization.py`): enumerate the families and record, per family, the child, parents, flattened atom scope and the LCN sentences *attached* to that family.
2. **Local credal sets** (`cn/local_credal_sets.py`): for each family and each interpretation of its scope, solve a (fractional) optimization problem to obtain the lower/upper bound $[p, \bar{p}]$ of the conditional $P(\text{ch}_v = c \mid \text{pa}_v = \pi)$. The method axis carries exactly two settings (Section 1.1): `method="linear"` (the baseline) treats each family *in isolation*, and `method="linear-tight"` solves the same family against every in-scope LCN sentence plus the scope-restricted LMC equalities (scheme D1 of Section 4).

3. **Inference over the strong extension** (`cn/vertices.py`, `cn/cve.py`): enumerate the extreme points of each local credal set and run Credal Variable Elimination, which optimizes over the strong extension—each local credal set, for each parent configuration, contributes a freely and independently chosen extreme point.

1.1 Build API: two methods, two orthogonal knobs

The schemes below are exposed through the existing build entry points (`CredalNetwork.from_lcn`, `CredalNetworkVertices.from_lcn`) without multiplying flags. The previous `method="nlp"` option (which added only *pairwise* parent marginal independence) is **removed**: it is a strict half-measure for source A, dominated by D1 below, and its bilinear program added solver fragility for little tightening. The surface is therefore:

- `method` $\in$ {"linear", "linear-tight"}. "linear" is today's per-family-in-isolation baseline (unchanged). "linear-tight" is scheme D1: each family is solved against every LCN sentence whose atom scope lies inside the family scope, plus the scope-restricted LMC bilinear equalities.
- `solver` $\in$ {"ipopt", "scip"} (already present). On "linear-tight", `solver="scip"` is scheme D3: it solves the enriched local program to certified global optimality; "ipopt" is the fast, hardened (multi-restart + SLSQP-fallback) variant of the same program.
- `merge_budget` k (new): scheme D2. A pre-factorization pass merges adjacent families whose combined scope is $\leq k$ before `ChainGraphFactorization.build`; $k=1$ is no merge (today), $k = \text{full scope}$ is exact. Orthogonal to `method` and `solver`.

Thus D1 is the *core* of "linear-tight", while D2 (k) and D3 (`solver`) are independent cost-vs-tightness dials rather than separate method names. (This `cn/` factorization `method` is unrelated to the identically named parameters of `ibp.py` and `map/exact_map.py`.)

2 Running example: the alarm network

We use `examples/alarm.lcn` throughout. It has five atoms $\{A, B, C, D, E\}$ and five sentences:

```
s1: 0.1  <= P(B)                <= 0.2
s2: 0.05 <= P(E)                <= 0.1
s3: 0.6  <= P(A | (B or E))     <= 0.9
s4: 0.5  <= P(not(C xor D) | A) <= 0.7   ; True
s5: 0.05 <= P(A)                <= 0.1
```

Its structure is a chain graph. The simplified structure graph has the compound node $C-D$ (because $s4$ ties C and D together through $C \oplus D$), so the families and the LMC independencies the pipeline derives are:

Family	Attached sentences	LMC independencies
$P(A \mid B, E)$	$s1, s2, s3, s5$	$(B \perp E)$
$P(B)$	$s1$	$(C \perp B, E \mid D, A)$
$P(E)$	$s2$	$(D \perp B, E \mid C, A)$
$P(C-D \mid A)$	$s4, s5$	

The factorization is $P(A \mid B, E) P(C-D \mid A) P(B) P(E)$.

Local credal sets computed today (method="linear"). Solving each family in isolation gives the conditionals below (lower/upper, rounded). Note the many vacuous $[0, 1]$ entries—these are exactly where gap A bites.

$P(A \mid B, E)$		$P(C - D \mid A)$	
$A=0, B=0, E=0$	$[0.956, 1.0]$	$C=0, D=0, A=1$	$[0, 0.7]$
$A=1, B=0, E=0$	$[0, 0.044]$	$C=0, D=1, A=1$	$[0, 0.5]$
$A=0, B=0, E=1$	$[0, 1.0]$	$C=1, D=0, A=1$	$[0, 0.5]$
$A=1, B=0, E=1$	$[0, 1.0]$	$C=1, D=1, A=1$	$[0, 0.7]$
$\dots$ (rows with $B=1$ or $E=1$)	$[0, 1.0]$	$\dots$ ($A=0$ rows)	$[0, 1.0]$

Two phenomena stand out and recur in the discussion below:

- **s3 only constrains the $(B \vee E)$ corner.** s3 bounds $P(A \mid B \vee E)$, i.e. the conditional given the *disjunction*. It says nothing directly about $P(A \mid B=0, E=1)$ or $P(A \mid B=1, E=0)$ individually, so those parent configurations come back vacuous $[0, 1]$ —even though, jointly with s5 ($P(A) \in [0.05, 0.1]$) and the marginals s1, s2, the LCN pins them far tighter.
- **s4 constrains a sum, not a single state.** s4 bounds $P(\neg(C \oplus D) \mid A) = P(C=D \mid A)$, i.e. the sum of the two diagonal states $P(C=0, D=0 \mid A) + P(C=1, D=1 \mid A)$. The isolated family LP can therefore push any single state to its extreme (hence $[0, 0.7] / [0, 0.5]$), and for $A=0$ there is no constraint at all, giving $[0, 1]$.

This is the concrete object the schemes of Section 4 tighten.

3 Two sources of looseness

Write $\tilde{\mathcal{D}}$ for the joint credal set actually optimized by CVE. We have $\mathcal{D}_{\text{LCN}} \subseteq \tilde{\mathcal{D}}$, and the inclusion is typically strict for *two independent reasons*.

3.1 (A) Local-set widening

For family v the pipeline computes

$$K_v^{\text{local}} = \left\{ P(\text{ch}_v \mid \text{pa}_v) : \text{the sentences attached to family } v \text{ hold} \right\},$$

because `_linear_rows_and_obj` loads only `family["sentences"]`. The true projection of the LCN onto that family is

$$K_v^{\text{true}} = \left\{ P(\text{ch}_v \mid \text{pa}_v) : \exists P \in \mathcal{D}_{\text{LCN}} \text{ inducing this conditional} \right\},$$

and $K_v^{\text{true}} \subseteq K_v^{\text{local}}$, often strictly: sentences attached to *other* families and the LMC equalities further constrain the conditional but are discarded. `method="nlp"` repairs only one slice of this gap (pairwise parent marginal independence) and is still loose for ≥ 3 parents and for any constraint that spans families.

On the example. The $[0, 1]$ entries of $P(A \mid B=0, E=1)$ and of $P(C - D \mid A=0)$ are pure gap-A artifacts: s5 ($P(A) \leq 0.1$) and the parent marginals s1, s2 are not seen when family A is solved against only its attached rows, and family $C - D$ never sees that A is rare. The true projections are far narrower.

3.2 (B) Strong-extension widening

Even with $K_v = K_v^{\text{true}}$ exactly, the strong extension

$$\mathcal{D}_{\text{SE}} = \left\{ \prod_v P_v : P_v \in K_v \text{ chosen independently per parent-config} \right\}$$

allows every (node, parent-configuration) slot to pick its extreme point freely. The LCN’s non-local sentences *couple* choices across families that $\mathcal{D}_{\text{SE}}$ decouples, so $\mathcal{D}_{\text{LCN}} \subsetneq \mathcal{D}_{\text{SE}}$ in general. This is the credal set whose product structure `verify_lmc_factorization.py` already reconstructs (its `_JointReconstructor` takes one extreme point per (node, parent-config) slot and forms the product joint).

On the example. Gap B shows up in the $C-D$ family. The strong extension lets the four states of $P(C-D \mid A=1)$ be chosen so that $P(C=0, D=0 \mid A=1)$ takes its max 0.7 and $P(C=1, D=1 \mid A=1)$ takes its max 0.7 in the same joint—a vertex combination whose diagonal sum is 1.4, violating $s4$ ’s upper bound of 0.7. The LCN forbids that distribution; the strong extension admits it. No tightening of the *intervals* removes it, because the constraint couples two states of the same factor.

Either gap alone suffices to make posteriors loose; a tight scheme must close the one that dominates on the instance at hand.

4 Proposed schemes

D1 — Constraint back-propagation into local sets (attacks A; the new "linear-tight" method). When solving family v , replace the constraint source from “sentences attached to v ” with “every LCN sentence s whose atom scope satisfies $\text{atoms}(s) \subseteq \text{scope}(v)$,” and additionally append the LMC bilinear equalities (`lmc_constraint_groups.vec` in `utils/common.py`) restricted to $\text{scope}(v)$. This is exactly `method="linear-tight"`: when the enriched constraint set is purely linear the family stays an LP / Charnes–Cooper fractional LP, and when it acquires LMC bilinear equalities it is solved on the existing hardened ipopt (or, under `solver="scip"`, certified-global) path that the removed `"nlp"` option used to drive. Closes the part of gap A expressible within a single family scope; cannot capture constraints spanning atoms in different families. Near-zero structural change—only `_linear_rows_and_obj`’s sentence selection and a scope-restricted LMC row builder.

On the example. Family A ’s scope is $\{A, B, E\}$, so $s1, s2, s5$ all satisfy $\text{atoms}(s) \subseteq \text{scope}$ and would now be loaded together with $s3$. The disjunctive $s3$ plus $s5$ ($P(A) \leq 0.1$) plus the parent marginals collapse the vacuous $[0, 1]$ rows of $P(A \mid B, E)$ to genuine sub-intervals. Family $C-D$ now sees $s5$, tightening the $A=0$ rows. The diagonal-sum coupling of $s4$ is *within* the $C-D$ family scope, so this single change also removes the gap-B vertex combination of Section 3 from that local set.

D2 — Scope merging with a tightness dial (attacks A, partially B; the `merge_budget` knob). Introduce a budget k (`merge_budget`). Merge adjacent families whose combined scope size is $\leq k$ into a single *super-family* and compute one joint local credal set over the merged scope (a larger fractional LP / NLP). At $k = \text{full scope}$ this is exact inference; at $k = \text{single-family}$ it is today’s behavior. Provides an anytime tightness-vs-cost knob, with per-super-family cost 2^k . Implemented as a merging pass over `lcn.families` before `ChainGraphFactorization.build`, reusing the existing per-family solver on the larger scope; `estimate_difficulty` can choose k adaptively.

On the example. Merging families A and $C-D$ (they share the edge $A \rightarrow C-D$) into a super-family over $\{A, B, E, C, D\}$ makes $s3, s4, s5$ act jointly: the bound on $P(C=D \mid A)$ is now solved in a

model that also knows $P(A) \leq 0.1$ and the $A \mid B, E$ structure, so the $C-D$ conditionals tighten beyond what D1 achieves in isolation. Here $k=5$ is the full scope, i.e. exact; a budget $k=3$ would merge only where cheap.

D3 — Joint-LP certification of each local bound (attacks A; exact, certified; the solver="scip" path). For each (family, interpretation, sense), solve the bound for the family conditional $P(\text{ch}_v \mid \text{pa}_v)$ against the *full* LCN feasible region: all LCN sentences and all LMC equalities over the 2^n joint, with a fractional objective local to the family. This yields K_v^{true} exactly. It is reached by combining `method="linear-tight"` with `solver="scip"` (the existing global path) so the bounds carry a global optimality certificate. Cost is 2^n variables per solve; viable for small n or, more usefully, as the *certifier* of the cheaper schemes (it is essentially `ExactInference` run per conditional).

D4 — Region-restricted (extensive) credal network (attacks B; structural). Replace the free per-parent-config vertex choice with a *coupled* one. Variants: (4a) detect sentences that couple multiple families and add constraints in `cve.py`'s bucket elimination forbidding vertex combinations that violate them; (4b) introduce a shared latent "regime" parent to each coupled family so the local credal sets become conditional on a common selector, restoring the cross-family correlation the LCN imposes. Touches the CVE elimination, not the local solves.

On the example. D4 directly kills the *s4* violation of Section 3: the four extreme points of $P(C-D \mid A=1)$ may no longer be combined into a joint whose diagonal sum exceeds 0.7. Under variant 4a, CVE's bucket for $C-D$ carries the side constraint $P(C=0, D=0 \mid A=1) + P(C=1, D=1 \mid A=1) \leq 0.7$ across the vertex choice, restoring the coupling that the free product had dropped.

Implementation note (what D4 can and cannot couple). The *s4* example above is, on inspection, an *in-family* constraint ($\text{atoms}(s4) = \{C, D, A\} \subseteq$ the $C-D$ family scope), so D1 already removes it; it is shown only to illustrate the coupling mechanism. The constraints D4 must actually handle are the *cross-family* ones — those whose atom set is a subset of no single family scope. A subtle but important fact constrains the menu: the strong extension of a chain-graph credal network is a *product* $\prod_v P(\text{ch}_v \mid \text{pa}_v)$, and that product structure *satisfies every LMC independence by construction* — which is exactly why `verify_lmc_factorization.py` reports the factorization SUPPORTED on every chain LCN we have tried. Consequently *no* product joint can violate a cross-family *LMC assertion*: enforcing them in D4 is sound but inert. The gap-B looseness D4 removes is therefore carried by cross-family *sentences* (Type-1/Type-2 bounds whose scope spans families), which the free product *can* violate because each local factor is chosen to satisfy only its own attached sentences. As a concrete witness, a three-family instance $P(A), P(B \mid A), P(C \mid A)$ with the cross-family sentence $P(B \wedge C) \leq 0.05$ admits a free-product joint reaching $P(B \wedge C) = 0.92$; D4 forbids it, collapsing e.g. $P(B \mid C=1)$ from $[0.04, 0.96]$ to a near-point. The implementation (`cn/coupling.py`, the `coupling="cross-family"` flag on Credal VE / CTE / ApproxLP) classifies a constraint as cross-family against the *built* factors, so a D2 `merge_budget` that fuses the coupled families turns the same constraint in-family and removes it from D4's list — D2 and D4 compose without double-counting.

D5 — Constraint-aware CVE (attacks A and B; the tight target). Keep the chain-graph structure for the elimination order and bucket scopes, but at each bucket combination optimize the bucket potential subject to the original LCN sentences *projected onto the bucket scope*, rather than over the free strong-extension vertices. This is a sequence of small NLPs threaded through variable elimination—a credal VE that respects the global constraints at every step and never materializes

the loose strong extension. In the limit it equals exact inference but its cost is bounded by the largest bucket (treewidth), not 2^n .

On the example. Eliminating B, E first produces a bucket over $\{A\}$ whose potential is optimized subject to s_1, s_2, s_3, s_5 jointly (closing gap A for A); eliminating C, D uses a bucket over $\{A\}$ optimized subject to s_4 as a coupled sum (closing gap B for $C-D$). The query marginal $P(A)$ is then read off the final $\{A\}$ bucket. No step ever forms the loose product, and no bucket exceeds three atoms.

Implementation note (junction-tree exact NLP). D5 is realized (`cn/junction_nlp.py`) as its own **CredalJT** engine — a build-once junction tree shared across all marginals: the clique tree and the constraint NLP (cluster variables, separator equalities, sentence and LMC rows) are query-independent, so they are constructed a single time and only the objective $P(\text{atom}=1 \mid e)$ is swapped per atom (two NLP solves each). (D5 used to live behind **CredalVE**’s `coupling="d5"`, which rebuilt the entire tree and model once *per atom*; that option has been removed and Credal VE now covers only `off/cross-family`.) The tree is built not as a sequence of message NLPs but as a single equivalent *junction-tree* NLP, which is cleaner to make provably exact. Build a clique tree from the chain-graph elimination order (reusing the bucket-tree construction), introduce one variable vector per cluster ranging over that cluster’s joint distribution, tie adjacent clusters with Lauritzen separator-consistency equalities (the marginal of a cluster onto a separator equals the neighbour’s), and impose each LCN sentence and LMC equality on a cluster containing its atoms. Because the clique tree has the running-intersection property, a separator-consistent constraint-satisfying family of cluster marginals projects to exactly the marginals of some global joint in the LCN’s set, and conversely; hence min / max of the query marginal over this NLP equals the exact bound, at cost $\sum_c 2^{|\text{atoms}(c)|}$ rather than 2^n . Cross-family constraints (whose atoms span cliques) are handled by augmenting the elimination order with their node sets so they acquire a single host cluster — the JT analogue of D2’s family merging and D4’s order augmentation; the cluster a constraint forces sets the local treewidth, so on instances whose cross-family constraints are themselves near-global (e.g. **alarm**, whose LMC assertions $C \perp B, E \mid D, A$ span all five atoms) D5 degenerates to near-full-joint cost, while on instances with local cross-family coupling it stays cheap.

Two findings from the implementation. (1) The cluster NLP is nonconvex (Type-2 ratios, the bilinear LMC equalities, and the conditional-query ratio objective), so a local solver (ipopt) only reaches a local optimum that can fall *short* of the true bound; D5 must use the certified global backend (SCIP) to actually be exact, and defaults to it. (2) D5 is exact for *conditional* queries $P(q \mid e)$, where the vertex-enumeration methods (plain Credal VE, and D4 on top of it) are *not even sound*: the conditional bound is a ratio whose extremum need not lie at a vertex of the unconditioned credal set, so min/max-over-vertices can land strictly inside the true interval. On **examples/d4.biting.lcn**, $P(B \mid C=1)$ is exactly $[0.006, 0.926]$ (D5, matching the global solver), whereas plain Credal VE reports $[0.038, 0.963]$ and D4 reports the spurious near-point $[0.038, 0.038]$. D5 is therefore the only member of the family that is sound for conditional queries, not merely tighter.

5 Detailed algorithms and worked examples

This section gives a step-by-step description of each scheme as implemented in the `cn/` package, together with a fully worked example on a running instance and a figure that traces the algorithm’s state. Two running instances are used:

- **alarm** (Section 2): five atoms $\{A, B, C, D, E\}$, the compound node $C-D$, families $P(A \mid B, E)$, $P(B)$, $P(E)$, $P(C-D \mid A)$. Used for D1, D2, D3 (local-set tightening, gap A).

Scheme	Source closed	Tightness	Cost	Effort
D1 back-propagate in-scope constraints	A (partial)	+	≈ 0	low
D2 scope merging (dial k)	A (+ B)	dial	2^k	low-med
D3 joint-LP certified locals	A (exact)	++	$2^n/\text{solve}$	med
D4 region-restricted CVE	B	+	med	med
D5 constraint-aware CVE	A + B	+++ ($\rightarrow$ exact)	treewidth	high

Table 1: Summary of the proposed schemes against the two looseness sources of Section 3.

$P(A \mid B, E)$ row	linear	linear-tight (D1)
$A=0, B=0, E=0$	[0.956, 1.0]	[0.988, 1.0]
$A=1, B=0, E=0$	[0, 0.044]	[0, 0.015]
$A=0, B=1, E=0$	[0, 1.0]	[0, 0.720]
$A=1, B=1, E=0$	[0, 1.0]	[0.280, 1.0]

Table 2: D1 on `alarm`’s A family: the vacuous $[0, 1]$ rows tighten once the scope-restricted constraints (here $s1, s2, s5$ and the $(B \perp E)$ equality) are imposed on the family joint. Bounds rounded; computed with `method="linear-tight"`.

- `d4_biting` (`examples/d4_biting.lcn`): three atoms, families $P(A)$, $P(B \mid A)$, $P(C \mid A)$ and the *cross-family* sentence $P(B \wedge C) \leq 0.05$. Used for D4 and D5 (cross-family coupling, gap B), because `alarm` has no cross-family *sentence* and so leaves gap B inert.

Throughout, an interpretation of a scope $X_1, \dots, X_k$ is packed *MSB-first*: the integer index of $(x_1, \dots, x_k)$ is $\sum_i x_i 2^{k-1-i}$. A local credal set for a factor $P(\text{ch} \mid \text{pa})$ is represented by its interval bounds $[\underline{p}, \bar{p}]$ per (child-state, parent-config); its *extreme points* are the finite set of conditional vectors realising those bounds, enumerated by the credal-net compiler (`cnv.extreme_points`).

5.1 D1 — constraint back-propagation (`method="linear-tight"`)

What it computes. For one family v with flattened atom scope $\text{sc}(v) = \text{ch}_v \cup \text{pa}_v$ and one interpretation π of that scope, D1 returns the lower/upper bound of the conditional $P(\text{ch}_v=c \mid \text{pa}_v=p)$ optimised not over the family’s attached sentences alone but over *all* sentences whose atoms lie in $\text{sc}(v)$ *and* the LMC independence equalities whose atoms lie in $\text{sc}(v)$. The pipeline already attaches every in-scope *sentence* to a family (via `process_chain_graph`), so the only delta over `method="linear"` is the addition of the scope-restricted *LMC bilinear equalities*; when a family has no in-scope LMC assertion, D1 falls back to the plain linear program and equals `"linear"`.

Worked example (`alarm`, family $P(A \mid B, E)$). The scope is $\{A, B, E\}$. The in-scope sentences are $s1 : P(B) \in [0.1, 0.2]$, $s2 : P(E) \in [0.05, 0.1]$, $s3 : P(A \mid B \vee E) \in [0.6, 0.9]$, $s5 : P(A) \in [0.05, 0.1]$. The one in-scope LMC assertion is $(B \perp E)$, contributing the *marginal* equality $P(B=1, E=1) = P(B=1)P(E=1)$ over the $\{A, B, E\}$ joint. Table 2 shows the effect: the rows that were vacuous $[0, 1]$ under `"linear"` (because $s3$ only constrains the $B \vee E$ corner) collapse to genuine sub-intervals once $(B \perp E)$ plus $s1, s2, s5$ act jointly on the family’s own joint. Figure 1 sketches the model.

Algorithm 1 D1 local credal set for one family (`_build_linear_tight_model`)

Require: family scope $sc = [s_1, \dots, s_k]$; attached sentences S ; LCN independencies $\mathcal{I}$; child state c , parent config p ; sense $\in \{\min, \max\}$

- 1: $\mathcal{A} \leftarrow \{a \in \mathcal{I} : \text{atoms}(a) \subseteq sc\}$ $\triangleright$ in-scope LMC assertions
- 2: **if** $\mathcal{A} = \emptyset$ **then return** `SOLVELINEAR`($sc, S, c, p, \text{sense}$) $\triangleright$ pure LP / Charnes–Cooper, identical to "linear"
- 3: **end if**
- 4: enumerate the $N = 2^k$ interpretations of sc ; create variables $q_0, \dots, q_{N-1} \geq 0$ $\triangleright$ a joint over the family scope
- 5: add the simplex constraint $\sum_i q_i = 1$
- 6: **for all** sentences $s \in S$ **do**
- 7: **if** s is Type-1 $\ell \leq P(\varphi) \leq u$ **then** add $\ell \leq \mathbf{1}_\varphi^\top q \leq u$
- 8: **else** (s is Type-2 $\ell \leq P(\varphi \mid \psi) \leq u$) add $\mathbf{1}_{\varphi\psi}^\top q \geq \ell \mathbf{1}_\psi^\top q$ and $\mathbf{1}_{\varphi\psi}^\top q \leq u \mathbf{1}_\psi^\top q$
- 9: **end if**
- 10: **end for**
- 11: **for all** $a = (X \perp Y \mid Z) \in \mathcal{A}$ and each group of `LMCGROUPS`(a, sc) **do**
- 12: **if** conditional group $(\mathbf{1}_{xyz}, \mathbf{1}_z, \mathbf{1}_{xz}, \mathbf{1}_{yz})$ **then** add $(\mathbf{1}_{xyz}^\top q)(\mathbf{1}_z^\top q) = (\mathbf{1}_{xz}^\top q)(\mathbf{1}_{yz}^\top q)$ $\triangleright$ bilinear
- 13: **else** (marginal group) add $\mathbf{1}_{xy}^\top q = (\mathbf{1}_x^\top q)(\mathbf{1}_y^\top q)$
- 14: **end if**
- 15: **end for**
- 16: build the objective $P(\text{ch}=c \mid \text{pa}=p) = (\mathbf{1}_{c,p}^\top q) / (\mathbf{1}_p^\top q)$ via the aux variable $t(\mathbf{1}_p^\top q) = \mathbf{1}_{c,p}^\top q$, $t \in [0, 1]$
- 17: **return** `OPTIMISE`(sense, t) with hardened ipopt (multi-restart) or, under `solver="scip"` (this is **D3**), the certified global solver

5.2 D2 — scope merging with budget k (`merge_budget`)

What it computes. D2 is a *pre-factorization* pass: before the local credal sets are solved, it merges adjacent families into super-families whose flattened scope is at most k atoms, then runs D1 (or "linear") on the larger scopes. Because a merged scope is a superset of each member scope, more sentences and more LMC assertions become in-scope, so the local sets can only tighten. At $k=1$ no merge happens (identical to today); at $k \geq n$ everything collapses into one family (exact). D2 operates on a *copy* of `lcn.families` and never mutates the model.

Worked example (alarm, $k=5$). The families are $P(A \mid B, E)$ (scope $\{A, B, E\}$), $P(B)$, $P(E)$, $P(C - D \mid A)$ (scope $\{C, D, A\}$). With $k=5$ the greedy pass contracts the arc $A \rightarrow (C - D)$ (their union $\{A, B, E, C, D\}$ has size $5 \leq k$), then folds B and E into the same super-family along the arcs $B \rightarrow A$, $E \rightarrow A$, producing a single family over all five atoms (Figure 2). Now *every* sentence and *every* LMC assertion is in-scope, including the cross-family assertions $(C \perp B, E \mid D, A)$ and $(D \perp B, E \mid C, A)$ that no 3-atom family could host — so the $k=5$ local set is the exact joint and the $C - D$ conditionals tighten beyond what D1 achieves on the 3-atom family. A budget $k=3$ would leave the families unchanged here (every pairwise union already exceeds 3), recovering plain D1.

5.3 D3 — certified local bounds (`solver="scip"`)

What it computes. D3 is not a separate model but a *backend*: it solves the D1 (or D2) local program with SCIP's spatial branch-and-bound global solver instead of the local ipopt. The D1




Figure 1: D1 (Algorithm 1). The family-scope joint q is constrained by the in-scope sentences *and* the scope-restricted LMC equalities; the conditional bound is read off as a fractional objective. "linear" omits the blue LMC block.




Figure 2: D2 (Algorithm 2) on `alarm` with $k=5$: greedy arc contraction collapses the four families into one super-family over $\{A, B, C, D, E\}$, bringing the cross-family LMC assertions in-scope. At $k=1$ the left DAG is used unchanged.

program is nonconvex (Type-2 ratio constraints and the bilinear LMC equalities), so `ipopt` only returns a *locally* optimal bound, which can be strictly looser *or* (for a feasibility-style extreme) spuriously tighter than the true local bound. `SCIP` returns the bound with a global optimality certificate (or a bounded gap under a time limit). D3 thus serves two roles: it produces the exact local credal set K_v^{true} for a family scope, and it *certifies* the cheaper D1/`ipopt` bounds.

Worked example (`alarm`, certifying $P(C=D \mid A)$). On the diagonal state $P(C=0, D=0 \mid A=1)$ the isolated family LP gives $[0, 0.7]$ (only the sum $P(C=D \mid A) \leq 0.7$ of $s4$ binds). Running the same model under `solver="scip"` returns the identical interval *with a zero optimality gap*, certifying that 0.7 is the true family-projected upper bound and that the `ipopt` result was not a local artifact. The contract D3 enforces is the containment

$$[\underline{p}^*, \bar{p}^*]_{\text{D3/scip}} \subseteq [\underline{p}, \bar{p}]_{\text{D1/ipopt}} \subseteq [\underline{p}, \bar{p}]_{\text{linear}},$$

checked per (family, row): a violation of the left inclusion flags an unsound `ipopt` local optimum, the width gap on the right quantifies how much looseness D1 still leaves (gap A residual). Figure 3 shows the role.

Algorithm 2 D2 family merging (`_merge_families`)

Require: families $\mathcal{F} = \{(\text{ch}_i, \text{pa}_i, S_i)\}$; budget k

```
1: build the family DAG: a node per family child; an arc  $u \rightarrow v$  whenever  $\text{ch}_u$  is a parent of  $v$ 
2: repeat
3:   for all arcs  $(u, v)$  in deterministic (sorted) order do
4:      $\text{sc} \leftarrow \text{flatten}(\text{scope}(u) \cup \text{scope}(v))$ 
5:     if  $|\text{sc}| > k$  then continue
6:     end if
7:     form the super-family  $M$ : child atoms =  $\text{ch}_u \cup \text{ch}_v$  (node name "-".join(sorted(...)));
       parents = external parents only (drop parents produced inside  $M$ )
8:     rewire every family that had  $u$  or  $v$  as a parent to point at  $M$ 
9:     if the rewired family DAG is acyclic then accept the contraction; restart the scan
10:    else reject (keep searching) ▷ D2 never creates a cycle
11:    end if
12:  end for
13: until no contraction was accepted
14: for all resulting (super-)families  $M$  do
15:   recompute attached sentences  $S_M = \{s : \text{atoms}(s) \subseteq \text{scope}(M)\}$  ▷ can only grow
16: end for
17: return the merged family list (fed to D1/"linear" on the larger scopes)
```

Algorithm 3 D3 certified local bound (`solver="scip"` on the D1 model)

Require: the D1 model $\mathcal{M}$ for (family, row, sense) from Algorithm 1

```
1: add the denominator floor  $\mathbf{1}_p^\top q \geq \varepsilon$  (so the conditional ratio stays well defined and SCIP cannot
   drive  $P(\text{pa}=p) \rightarrow 0$ )
2:  $(\text{val}, \text{gap}, \text{status}) \leftarrow \text{SCIPSOLVE}(\mathcal{M}, \text{time\_limit}, \text{gap\_tol})$ 
3: if  $\text{status} \in \{\text{OPTIMAL}\}$  or  $\text{gap} \leq \text{gap\_tol}$  then return certified val
4: else return best-so-far val with the reported gap ▷ partial certificate
5: end if
```

5.4 D4 — region-restricted coupling (`coupling="cross-family"`)

What it computes. Credal VE represents each node’s local credal set as a **Potential**: a finite set of *functions* (numpy arrays), each one a fully instantiated CPT obtained by picking, independently per parent-configuration, one local extreme point. Bucket elimination repeatedly *combines* potentials (a Cartesian product of their function sets) and *marginalises* a variable out; the final query bound is the min / max of the normalised query state over the surviving functions. This is the free strong-extension product. D4 *drops* any function whose assembled joint violates a *cross-family* constraint — a sentence or LMC assertion whose atom set is a subset of no single family scope. Crucially, a fully assembled function over a node set *is* the unnormalised joint over those nodes’ atoms, so a cross-family residual can be evaluated on it directly, reusing exactly the verifier’s residual math.

The two pieces that make it correct.

1. **Checkability gate.** A constraint is evaluated on a function only once the function’s scope covers *all* the constraint’s atoms; until then the predicate returns “feasible” (do not drop).




Figure 3: D3 (Algorithm 3) solves the D1 model with the global solver, certifying the local bound. The nesting $D3 \subseteq D1 \subseteq \text{linear}$ is the per-row validation contract (Section 6).

Algorithm 4 D4 coupled Credal VE (coupling="cross-family")

Require: credal net cnv ; query q ; evidence e

- 1: $\mathcal{C} \leftarrow$ cross-family constraints $\{c : \text{atoms}(c) \not\subseteq \text{scope}(f) \ \forall \text{ families } f\}$
 - 2: **if** $\mathcal{C} = \emptyset$ **then** run plain Credal VE $\triangleright$ D4 inert; e.g. **alarm**
 - 3: **end if**
 - 4: $order \leftarrow \text{MINFILL}(\text{factor scopes} \cup \{\text{node set of } c : c \in \mathcal{C}\}, \text{exclude}=\{q\})$
 - 5: build one **Potential** per node from its extreme points; add evidence indicators
 - 6: **for all** v in $order$ **do**
 - 7: $bucket \leftarrow$ potentials whose scope contains v ; $combined \leftarrow$ product of the bucket
 - 8: **D4 filter:** keep $f \in combined$ iff $\text{FEASIBLE}(f, \text{scope}(combined), \mathcal{C})$ $\triangleright$ on the full bucket scope, before marginalising
 - 9: $result \leftarrow \text{MARGINALISE}(combined, v)$; $result \leftarrow \text{PRUNE}(result)$
 - 10: **end for**
 - 11: combine the leftovers; **return** min / max of the normalised q state over the surviving functions

 - 12: **function** $\text{FEASIBLE}(f, sc, \mathcal{C})$
 - 13: **for all** $c \in \mathcal{C}$ with $\text{atoms}(c) \subseteq sc$ **do**
 - 14: marginalise f to $\text{atoms}(c)$, normalise, evaluate the residual (Type-1/2 bound or bilinear LMC); **if** violated **return** false
 - 15: **end for**
 - 16: **return** true $\triangleright$ atoms not yet all in scope $\Rightarrow$ not droppable
 - 17: **end function**
-

This makes filtering after *any* elimination step sound: it can only ever remove a genuinely infeasible function, never one whose feasibility is undecided.

2. **Order augmentation.** For the gate to ever fire, the constrained atoms must co-occur in some bucket. D4 therefore forces the min-fill elimination order, augmented with one extra “clique scope” per cross-family constraint (its node set), so those nodes land in a common bucket — otherwise an atom is summed out before the constraint scope assembles and D4 stays inert (sound but useless).

Worked example (d4.biting, query $P(B \mid C=1)$). The cross-family sentence is $P(B \wedge C) \leq 0.05$ (atoms $\{B, C\}$, fitting neither the B nor the C family). The free product can pick $P(B=1 \mid A)$ and $P(C=1 \mid A)$ both large, giving a joint with $P(B \wedge C) = 0.92$ — admissible in the strong extension but forbidden by the LCN. Augmenting the order forces $\{A, B, C\}$ into one bucket; the D4 filter there drops every function whose $P(B \wedge C)$ exceeds 0.05. Figure 4 traces it. Net effect on the *unconditional* marginals is small, but on a query that makes B and C interact (e.g. evidence $C=1$) the upper bound on $P(B \mid C=1)$ drops sharply.




Figure 4: D4 (Algorithm 4) on `d4.biting`: in the augmented bucket $\{A, B, C\}$, functions whose joint violates the cross-family $P(B \wedge C) \leq 0.05$ (e.g. f_3 with 0.92) are dropped *before* A is summed out; the survivors yield the tightened query.

A sharp limitation (conditional queries). D4 is sound only as a tightening of the strong-extension bound, and the strong-extension bound is itself a valid *outer* bound only for *unconditional* marginals. For a conditional query $P(q \mid e)$ the bound is a ratio whose extremum need not lie at a vertex of the unconditioned credal set, so min / max-over-vertices — and hence D4’s filtered version — can land *strictly inside* the true interval. On `d4.biting`, $P(B \mid C=1)$ is truly $[0.006, 0.926]$, yet plain Credal VE reports $[0.038, 0.963]$ and D4 the spurious near-point $[0.038, 0.038]$. The implementation prints a warning when `coupling="cross-family"` is used with evidence and directs the user to D5 (Section 5.5), which is exact.

5.5 D5 — junction-tree exact NLP (the CredalJT engine)

What it computes. D5 abandons the vertex product entirely and computes the *exact* chain-graph marginal by a single nonlinear program over *cluster-marginal* variables. Build a clique (junction) tree from the chain-graph elimination order; give each cluster c a variable vector q_c ranging over the joint distribution of its atoms; tie adjacent clusters with Lauritzen *separator-consistency* equalities (the marginal of one cluster onto a shared separator equals the neighbour’s); and impose every LCN sentence and LMC equality on a cluster whose atom set contains it. Reading the query marginal off the cluster that holds the query (and evidence) and optimising it min / max over this feasible region returns the exact LCN bound. Because each cluster is bounded by the treewidth, the NLP has $\sum_c 2^{|\text{atoms}(c)|}$ variables rather than 2^n .

Why it is exact. For a chain graph the clique tree has the *running intersection property* (RIP). With RIP, a family of cluster marginals that (i) agrees on every separator and (ii) satisfies, in some containing cluster, every LCN/LMC constraint, extends to a single global joint feasible for the full 2^n **ExactInference** model — and conversely every feasible global joint projects to such a family. Hence the cluster NLP and the full-joint NLP have the same optimal value. Cross-family constraints (and the query+evidence atoms) are forced to share a host cluster by augmenting the elimination order with their node sets, exactly as in D4; if that would push a cluster past the treewidth budget, D5 falls back to **ExactInference.run_query** (still exact, at 2^n). This exactness argument is not specific to chains and is made precise (with the exact condition under which a *dropped* LMC is still respected, and the one structure where it fails) as Theorem 1 in Section 5.6.

Algorithm 5 D5 junction-tree exact NLP (**CredalJT.run**; one tree, all marginals)

Require: credal net cnv ; evidence e ; budget $k_{\max}$; solver

- 1: $\mathcal{C} \leftarrow$ cross-family constraints; augment the min-fill order with their node sets *only* (no query — the tree is query-INDEPENDENT)
 - 2: build the clique tree T (clusters $\mathcal{K}$, edges with separators) *once*; flatten each cluster to its atom set; prune leaf clusters subsumed by a neighbour
 - 3: **if** $\max_c |\text{atoms}(c)| > k_{\max}$ **then return** EXACTINFERENCE per atom $\triangleright$ budget fallback
 - 4: **end if**
 - 5: (*build the constraint model once*) create variables $q_c[i] \geq 0$ per cluster and state, with $\sum_i q_c[i] = 1$
 - 6: **for all** tree edges (c, d) with separator S **do**
 - 7: add the separator-consistency equalities $M_{c \rightarrow S} q_c = M_{d \rightarrow S} q_d$
 - 8: **end for**
 - 9: **for all** sentences / LMC assertions **do** add their row(s) on a host cluster containing their atoms (Type-1/2 linear/ratio, or bilinear LMC)
 - 10: **end for**
 - 11: **for all** non-evidence singleton atoms a **do** $\triangleright$ loop; only the objective changes
 - 12: pick a host cluster c_a containing a and the evidence atoms; build the objective $P(a=1 \mid e) = (\mathbf{1}_{a,e}^\top q_{c_a}) / (\mathbf{1}_e^\top q_{c_a})$ via the aux-variable trick (denominator floor under SCIP)
 - 13: $\underline{q}_a \leftarrow \text{SOLVE}(\min)$, $\bar{q}_a \leftarrow \text{SOLVE}(\max)$ over the *shared* model with the global backend (SCIP)
 - 14: **end for**
 - 15: **return** $\{a \mapsto [\underline{q}_a, \bar{q}_a]\}$
-

Worked example (d4_biting, query $P(B \mid C=1)$). The cross-family sentence $P(B \wedge C) \leq 0.05$ has node set $\{B, C\}$, and the query+evidence atoms are $\{B, C\}$; augmenting forces them together, so the elimination yields a clique tree with a cluster $c_q = \{A, B, C\}$ (the full joint) and a subsumed leaf that is pruned. All four sentences, the LMC assertion $(B \perp C \mid A)$, and the objective all host on c_q . The NLP is therefore identical to the full-joint **ExactInference** model, and solving it with SCIP returns $P(B \mid C=1) = [0.006, 0.926]$ — the exact bound, matching the certified global oracle and correcting both plain Credal VE $([0.038, 0.963])$ and D4 $([0.038, 0.038])$. Figure 5 shows the cluster tree and the NLP it induces.

Two implementation findings. (1) The cluster NLP is nonconvex (Type-2 ratios, bilinear LMC, the conditional-ratio objective); the local solver ipopt returns a local optimum that can fall *short* of the bound (it gave 0.90 instead of 0.926 above), so D5 defaults to the certified global backend SCIP. (2) The treewidth saving requires that the *junction tree* stay small. Only cross-family *sentences* are allowed to enlarge a cluster (by augmenting the elimination order); LMC assertions are not, because they are structural — a chain-graph junction tree enforces them through its separators (running intersection), so **CredalJT** skips any LMC that is already RIP-implied by the tree and drops the rest (a chain-graph product joint satisfies every LMC by construction). This is what lets a long chain run at treewidth (Section 5.5.2) and keeps **alarm** at treewidth 3 rather than collapsing to the full 2^5 joint, even though its assertions $(C \perp B, E \mid D, A)$, $(D \perp B, E \mid C, A)$ name all five atoms.




Figure 5: D5 (Algorithm 5) on `d4_biting`: the augmented clique tree has the full-joint cluster $c_q = \{A, B, C\}$ hosting all constraints and the conditional objective; separator-consistency equalities (right) tie clusters on general instances. Solving with SCIP gives the exact $P(B | C=1)$.

5.5.1 A complete worked example: all marginals of `d4_biting` with CredalJT

This subsection works the `CredalJT` engine end to end on `examples/d4_biting.lcn`, computing *all* singleton marginals with no evidence, and shows the junction tree, the messages, and exactly how each message is computed.

The model. `d4_biting` has three atoms $\{A, B, C\}$ and four sentences:

$$\begin{aligned}
 A1: & 0.4 \leq P(A) && \leq 0.6 \\
 B1: & 0.1 \leq P(B | A) && \leq 0.3 \\
 C1: & 0.2 \leq P(C | A) && \leq 0.5 \\
 BC: & 0.0 \leq P(B \text{ and } C) && \leq 0.05
 \end{aligned}$$

A is a root; B and C are its children. The families are $P(A)$, $P(B | A)$, $P(C | A)$, with the LMC independence $(B \perp C | A)$. The sentence BC (atoms $\{B, C\}$) is *cross-family*: it constrains B and C jointly but lies in no single family. This is what makes the strong extension loose—the free product can pick $P(B=1 | A)$ and $P(C=1 | A)$ together so that $P(B \wedge C)$ reaches 0.92, far above the 0.05 bound—and what D5 enforces exactly.

Local credal sets. The `cn/` compiler solves each family in isolation (`method="linear"`), giving the conditional intervals (lower/upper):

$P(A)$	$P(B A)$	$P(C A)$
$A=1 : [0.4, 0.6]$	$B=1, A=1 : [0.1, 0.3]$	$C=1, A=1 : [0.2, 0.5]$
	$B=1, A=0 : [0, 1]$	$C=1, A=0 : [0, 1]$

(B, C are unconstrained given $A=0$ because $B1, C1$ condition on A and $A=0$ is the complement; BC does not appear in any single family.)




Figure 6: The `d4_biting` junction tree built by `CredalJT` (one tree for all marginals). Cluster $A = \{A, B, C\}$ is the full joint and hosts every sentence and the LMC equality; the separators $\{B, C\}$ and $\{C\}$ tie the downstream clusters B and C to it.

The junction tree. `CredalJT` builds one *query-independent* clique tree (Algorithm 5). The cross-family sentence BC contributes the augmenting clique $\{B, C\}$, which—together with the family scopes $\{A\}$, $\{A, B\}$, $\{A, C\}$ and the moralization A couples B and C —makes the min-fill elimination produce three clusters and the chain

$$C : \{C\} \xleftarrow{S=\{C\}} B : \{B, C\} \xleftarrow{S=\{B, C\}} A : \{A, B, C\}.$$

Figure 6 draws it. Cluster A has the full atom set $\{A, B, C\}$ and hence is the host of *everything*: all four sentences *and* the LMC equality $(B \perp C \mid A)$ sit on A ; clusters B and C carry no constraints and exist only to expose the marginals $P(B, C)$ and $P(C)$ through the separators. (The pruning step keeps B and C because in the tree orientation $C \leftarrow B \leftarrow A$ each is the *parent* of the larger cluster, not a subsumed leaf; they are harmless—each merely copies a marginal of A .)

The NLP variables. One probability vector per cluster: $q_A \in \Delta(2^3)$ (8 states over A, B, C , MSB-first), $q_B \in \Delta(2^2)$ (4 states over B, C), $q_C \in \Delta(2^1)$ (2 states over C); each on its own simplex. Abbreviate the cells $a_i \equiv q_A[i]$, $b_j \equiv q_B[j]$, $c_k \equiv q_C[k]$ with the MSB-first decoding

$a_0 = P(A0B0C0)$	$a_1 = P(A0B0C1)$	$a_2 = P(A0B1C0)$	$a_3 = P(A0B1C1)$
$a_4 = P(A1B0C0)$	$a_5 = P(A1B0C1)$	$a_6 = P(A1B1C0)$	$a_7 = P(A1B1C1)$
$b_0 = P(B0C0)$	$b_1 = P(B0C1)$	$b_2 = P(B1C0)$	$b_3 = P(B1C1)$
$c_0 = P(C0)$	$c_1 = P(C1)$		

(“ $A1B0C1$ ” is shorthand for $A=1, B=0, C=1$.)

The NLP solved (exactly as `CredalJT` builds it). The shared constraint model is the following 14-variable program. The bracketed labels are the constraint kind. *Simplexes*:

$$\sum_{i=0}^7 a_i = 1, \quad \sum_{j=0}^3 b_j = 1, \quad c_0 + c_1 = 1. \quad [\text{simplex}]$$

Sentence rows on cluster A (here $P(A=1) = a_4 + a_5 + a_6 + a_7$, $P(B=1, A=1) = a_6 + a_7$, $P(C=1, A=1) = a_5 + a_7$, $P(B=1, C=1) = a_3 + a_7$):

$$\begin{aligned} 0.4 &\leq a_4 + a_5 + a_6 + a_7 \leq 0.6 && [A1] \\ 0.1(a_4 + a_5 + a_6 + a_7) &\leq a_6 + a_7 \leq 0.3(a_4 + a_5 + a_6 + a_7) && [B1 \text{ (Type-2 ratio)}] \\ 0.2(a_4 + a_5 + a_6 + a_7) &\leq a_5 + a_7 \leq 0.5(a_4 + a_5 + a_6 + a_7) && [C1 \text{ (Type-2 ratio)}] \\ 0 &\leq a_3 + a_7 \leq 0.05 && [BC \text{ (cross-family)}] \end{aligned}$$

Bilinear LMC equalities for $(B \perp C \mid A)$ — the factorization $P(b, c \mid A) = P(b \mid A) P(c \mid A)$ written as $P(b, c, A)P(A) = P(b, A)P(c, A)$ on each (b, c) slice of each value of A :

$$\begin{aligned} a_2(a_0+a_1+a_2+a_3) &= (a_2+a_3)(a_0+a_2), & a_3(a_0+a_1+a_2+a_3) &= (a_2+a_3)(a_1+a_3), & [A=0] \\ a_6(a_4+a_5+a_6+a_7) &= (a_6+a_7)(a_4+a_6), & a_7(a_4+a_5+a_6+a_7) &= (a_6+a_7)(a_5+a_7). & [A=1] \end{aligned}$$

Separator-consistency equalities ($A \rightarrow B$ over $\{B, C\}$, then $B \rightarrow C$ over $\{C\}$):

$$\begin{aligned} b_0 &= a_0 + a_4, & b_1 &= a_1 + a_5, & b_2 &= a_2 + a_6, & b_3 &= a_3 + a_7, & [A \rightarrow B] \\ c_0 &= b_0 + b_2, & c_1 &= b_1 + b_3. & & & & & [B \rightarrow C] \end{aligned}$$

Two observations make the structure transparent. First, the sentence and LMC rows live *only* on cluster A (the full joint); b and c are pinned purely by the separator equalities, so they carry no independent freedom—they are exactly the marginals $P(B, C)$ and $P(C)$ of a . Second, the cross-family row BC ($a_3+a_7 \leq 0.05$) is the only one coupling the B - and C -children; it is what the free product violates and the joint cluster A enforces.

What the separator equalities mean (the messages). The two “[$A \rightarrow B$]” and “[$B \rightarrow C$]” blocks above are the Lauritzen separator-consistency conditions, and they are precisely the junction-tree messages. The edge $A \rightarrow B$ forces each q_B cell to equal the sum of the q_A cells that agree on (B, C) — i.e. q_B must equal the marginal of the cluster- A joint onto the separator $\{B, C\}$; that marginal (b_0, b_1, b_2, b_3) is the “message” A sends to B . The edge $B \rightarrow C$ likewise makes (c_0, c_1) the marginal of q_B onto $\{C\}$. Because the clique tree has the running-intersection property, these two edge conditions are enough to guarantee a single global joint over $\{A, B, C\}$ consistent with all three cluster marginals — which is why optimizing on cluster A alone yields the exact bound.

Reading a marginal: the two NLPs for $P(B)$. For each non-evidence atom the engine appends the atom’s objective to the *shared* constraint model and solves it twice. For B , the objective is $P(B=1) = a_2 + a_3 + a_6 + a_7$ (the $B=1$ cells of q_A ; no evidence, so the objective is linear and no auxiliary ratio variable is needed). The two programs are

$$\underline{P(B=1)} = \min (a_2+a_3+a_6+a_7) \quad \text{and} \quad \overline{P(B=1)} = \max (a_2+a_3+a_6+a_7)$$

both **subject to the full constraint set above** (the three simplexes, the four sentence rows $A1, B1, C1, BC$, the four bilinear LMC equalities, and the six separator equalities). These are nonconvex (Type-2 ratios and the bilinear LMC), so they are solved to certified global optimality with SCIP. The maximization returns $\overline{P(B=1)} = 0.72$ at the cluster- A joint

state (A, B, C)	a_0	a_1	a_2	a_3	a_4	a_5	a_6	a_7
	000	001	010	011	100	101	110	111
q_A	0	0	0.5751	0.0249	0.2227	0.0573	0.0954	0.0246

One can check this point against the constraints: $P(A=1) = a_4+a_5+a_6+a_7 = 0.40$ (at $A1$ ’s lower bound); $B1$: $a_6+a_7 = 0.12 = 0.3 \cdot 0.40$ (at the upper bound, pushing B up); $C1$: $a_5+a_7 = 0.0819 \approx 0.2 \cdot 0.40$; and crucially the cross-family BC is *binding*, $a_3+a_7 = 0.0249+0.0246 = 0.0495 \leq 0.05$. The objective is $P(B=1) = 0.5751+0.0249+0.0954+0.0246 = 0.72$. The *message* this maximizer sends from A to B is its marginal onto $\{B, C\}$, computed cell-by-cell by the separator equalities:

$$(b_0, b_1, b_2, b_3) = (a_0+a_4, a_1+a_5, a_2+a_6, a_3+a_7) = (0.2227, 0.0573, 0.6705, 0.0495),$$

e.g. $b_2 = P(B=1, C=0) = a_2 + a_6 = 0.5751 + 0.0954 = 0.6705$; and onward $B \rightarrow C$, $(c_0, c_1) = (b_0 + b_2, b_1 + b_3) = (0.8932, 0.1068)$. Without the joint cluster (i.e. over the free product) B could be pushed to $P(B=1)$ near 1 while C stayed high; the 0.05 cap on the shared cell $a_3 + a_7$ is exactly what holds $\overline{P(B=1)}$ down to 0.72. The minimization, over the same constraint set, gives $\overline{P(B=1)} = 0.04$, so $P(B) \in [0.04, 0.72]$.

All marginals. Running every atom through the same model yields

$$P(A) = [0.40, 0.60], \quad P(B) = [0.04, 0.72], \quad P(C) = [0.08, 0.80],$$

each matching the certified-global **ExactInference** oracle to solver tolerance. The junction tree and constraint model are built *once*; only the objective changes across the three atoms (two SCIP solves each). At **verbosity=2** **CredalJT** prints exactly this information: the tree, the per-state separator-consistency equalities, and the realized message values at each min/max optimizer.

With evidence: the conditional NLP. The example above has no evidence, so each objective is the *linear* sum of the atom’s 1-cells. With evidence e , the objective is the conditional ratio $P(a=1 \mid e) = P(a=1, e)/P(e)$, which the engine linearizes with one auxiliary variable $t \in [0, 1]$ on the host cluster: it adds the bilinear equality $t \cdot (\mathbf{1}_e^\top q_{c_a}) = \mathbf{1}_{a,e}^\top q_{c_a}$ (and, under SCIP, the denominator floor $\mathbf{1}_e^\top q_{c_a} \geq \varepsilon$ so $P(e)$ cannot vanish), and optimizes min/max t . Everything else—the same shared constraint model and the same separator messages—is reused; only this small objective block is added and cleared between atoms.

5.5.2 The treewidth advantage: a Markov-chain example

The headline claim for D5 is that it operates in space exponential in the *treewidth*, not in the number of variables. This subsection makes that concrete on a clean chain (**examples/chain.lcn**) and shows how the engine avoids the full joint.

The instance. A binary Markov chain $x_0 \rightarrow x_1 \rightarrow \dots \rightarrow x_7$ ($n = 8$): a prior on x_0 and one transition per link,

$$\begin{aligned} P0 : 0.3 &\leq P(x_0) && \leq 0.5 \\ T_i : 0.2 &\leq P(x_{i+1} \mid x_i) && \leq 0.4 \quad \text{and} \quad 0.6 \leq P(x_{i+1} \mid \neg x_i) \leq 0.8 \end{aligned}$$

for $i = 0, \dots, 6$. *Every* sentence is family-local; there is no cross-family sentence. The LCN’s only cross-family constraints are the LMC assertions $(x_{i+1} \perp \{x_0, \dots, x_{i-1}\} \mid x_i)$, whose **all_vars** grow up to all eight atoms.

Why a naive reading would blow up — and why D5 does not. If those LMC assertions were forced into a single host cluster (as a literal “all atoms in one bucket” reading demands), the engine would build one $2^8 = 256$ -state cluster: no better than the full joint. The key is that **for a chain graph the LMC assertions are structural** — they are exactly the independencies the junction tree already encodes through its separators. **CredalJT** therefore (i) augments the elimination order with cross-family *sentence* node sets only (here there are none), and (ii) imposes an LMC assertion only if it is *not* already RIP-implied by the resulting tree. On the chain *every* LMC is RIP-implied ($Z = \{x_i\}$ is the separator on the path between x_{i+1} and the earlier atoms), so *none* are imposed, and the tree is the linear clique chain

$$\{x_0, x_1\} \overset{\{x_1\}}{-} \{x_1, x_2\} \overset{\{x_2\}}{-} \dots \overset{\{x_6\}}{-} \{x_6, x_7\},$$




Figure 7: **CredalJT** on the 8-atom chain `chain.lcn`: the junction tree is a linear chain of 7 two-atom clusters (treewidth 1) tied by single-atom separators. The exact joint over all eight atoms is $2^8 = 256$ states, but the engine never materializes it — it solves an NLP with $7 \cdot 2^2 = 28$ cluster-marginal variables.

i.e. 7 clusters of two atoms each plus a trivial $\{x_7\}$ leaf, **treewidth** 1 (Figure 7). The NLP has $\sum_c 2^{|\text{atoms}(c)|} = 7 \cdot 2^2 = 28$ cluster-marginal variables, versus $2^8 = 256$ for the full joint — and the gap is $n \cdot 2^2$ vs 2^n , i.e. *linear* in n instead of exponential.

Step by step. (1) Build the families $P(x_0), P(x_1 | x_0), \dots, P(x_7 | x_6)$ (scopes $\{x_i, x_{i+1}\}$). (2) Min-fill over the family scopes (no augmentation, since there is no cross-family sentence) yields the clique chain above; each x_i is a separator. (3) Check each LMC: $(x_{i+1} \perp x_{<i} | x_i)$ has $Z = \{x_i\} \subseteq$ the separator $\{x_i\}$, with x_{i+1} on one side and $x_{<i}$ on the other $\Rightarrow$ RIP-implied $\Rightarrow$ skipped. (4) Build one constraint NLP: per cluster c_i a vector $q_{c_i} \in \Delta(2^2)$ over (x_i, x_{i+1}) with a simplex; the link sentences $P(x_{i+1} | x_i)$ as Type-2 ratio rows on c_i ; and the separator equalities $M_{c_i \rightarrow \{x_{i+1}\}} q_{c_i} = M_{c_{i+1} \rightarrow \{x_{i+1}\}} q_{c_{i+1}}$. (5) For each atom, read its marginal off any hosting cluster and solve min / max over this shared model.

One message, concretely. The edge $c_0 \rightarrow c_1$ over separator $\{x_1\}$ enforces, for the solved distribution, q_{c_1} ’s x_1 -marginal = q_{c_0} ’s x_1 -marginal: $q_{c_0}[x_0=0, x_1=1] + q_{c_0}[x_0=1, x_1=1] = q_{c_1}[x_1=1, x_2=0] + q_{c_1}[x_1=1, x_2=1]$. That shared value of $P(x_1=1)$ is the message passed from the $\{x_0, x_1\}$ block into the $\{x_1, x_2\}$ block; it is all c_1 needs to know about x_0 .

Comparison with the other engines. All of the `cn/` engines consume the same credal network. On this chain there is no cross-family *sentence*, so the strong extension equals the true LCN set and every engine reports the same exact interval — but at very different cost and with different guarantees:

	$P(x_0=1)$	$P(x_2=1)$	$P(x_7=1)$	cost / character
ExactInference	[0.300, 0.500]	[0.328, 0.640]	[0.334, 0.667]	full joint, $2^8 = 256$ vars
CredalJT (D5)	[0.300, 0.500]	[0.328, 0.640]	[0.334, 0.667]	treewidth , $7 \cdot 2^2 = 28$ vars; exact
CredalVE	[0.300, 0.500]	[0.328, 0.640]	[0.334, 0.667]	per-target vertex elimination (strong ext.)
CredalCTE	[0.300, 0.500]	[0.328, 0.640]	[0.334, 0.667]	two-pass vertex tree (strong ext.)

CredalJT matches the full-joint **ExactInference** answer while working in treewidth-bounded space. **CredalVE**/**CredalCTE** bound the *strong extension*, which on a (sentence-free) chain coincides with the LCN set, so they too are exact here—but they enumerate and propagate the local extreme

vertices rather than exploiting the small treewidth of the NLP, and (unlike D5) would be merely outer bounds the moment a cross-family sentence is added. As n grows the contrast sharpens: D5 stays at $n \cdot 2^{w+1}$ (w the treewidth) while the full joint and any vertex enumeration that materializes it grow as 2^n .

Contrast with alarm. The favourable structure here is *local* cross-family coupling (none, in fact). **alarm** is the more interesting case: its cross-family LMC assertions $(C \perp B, E \mid D, A)$ and $(D \perp B, E \mid C, A)$ span all five atoms and are *not* single-edge-RIP-implied (the compound node $C - D$ hides D behind the $\{A\}$ separator, so the cheap `rip_implied` test misses them). **CredalJT** drops them, keeps the genuinely-binding small-scope $(B \perp E)$, and runs at treewidth 3. The drop is nonetheless *exactness-preserving* here, but for a subtler reason than “the product satisfies every LMC”: both assertions are *separator-entailed by the triangulated junction tree* — $\{A, D\}$ separates C from $\{B, E\}$ in the chordal graph G^* whose maximal cliques are the clusters $\{B, E, A\}, \{A, C, D\}$, so any separator-consistent family already satisfies them (Section 5.6 makes this precise). **CredalJT** therefore reproduces the brute-force-verified bounds $P(A)=[0.087, 0.100]$, $P(B)=[0.100, 0.123]$, $P(E)=[0.050, 0.074]$ — and is in fact *more* reliable than raw **ExactInference** on this instance, whose local solver over-tightens $P(A)$ to $[0.092, 0.100]$ and $P(E)$ to $[0.050, 0.073]$ (both *inside* the true interval; see the brute-force verifier in Section 5.6).

5.6 Generality and exactness of CredalJT

The Markov chain of Section 5.5.2 and the worked example of Section 5.5.1 might suggest that **CredalJT** is a chain-specific trick. It is not. This subsection answers two questions precisely: (i) *for which structures does the engine apply*, and (ii) *when is its output the exact LCN marginal bound* rather than merely a valid outer bound. Nothing in the construction privileges a chain; a chain is only the best-case *treewidth* ($w=1$). The same engine builds a genuine, branching junction tree on every chain-graph LCN, and is provably exact under an explicit, checkable condition. We state the condition, prove it, and report empirically (with an independent oracle) exactly where it holds and where it does not.

5.6.1 Setup and notation

Write $\mathbf{V} = \{V_1, \dots, V_n\}$ for the atoms. The LCN’s distribution set is the *full-joint feasible region* that **ExactInference** optimizes,

$$\mathcal{F} = \left\{ p \in \Delta(2^{\mathbf{V}}) : g_s(p) \in [\ell_s, u_s] \ \forall s, \quad h_\iota(p) = 0 \ \forall \iota \right\}, \quad (2)$$

where s ranges over the LCN sentences and ι over the LMC independence assertions. Each g_s is a Type-1 linear functional $P(\varphi) = \mathbf{1}_\varphi^\top p$ or a Type-2 ratio $P(\varphi \mid \psi)$ (encoded as the two bilinear inequalities $\mathbf{1}_{\varphi\psi}^\top p \geq b \mathbf{1}_\psi^\top p$); each h_ι is the bilinear LMC equality built by `lmc_constraint_groups_vec` for $(X \perp Y \mid S)$. The exact bound of a query atom a (optionally given evidence e) is $[\min, \max]$ of the read-off functional over $\mathcal{F}$, and this is precisely what **exact.py** computes on the 2^n joint.

The decisive structural fact, true by construction of the encoding and *independent of any graph topology*, is

Locality. *Every constraint depends on p only through its marginal on a fixed atom set, its scope: a sentence s through $p \upharpoonright \text{atoms}(s)$, and an LMC assertion $(X \perp Y \mid S)$ through $p \upharpoonright (X \cup Y \cup S)$.*

(A Type-1 row reads $\mathbf{1}_\varphi^\top p = \sum$ of marginal cells; a Type-2 row is a ratio of two such; an LMC group is a polynomial in marginal cells of $X \cup Y \cup S$. None ever inspects a cell outside its scope.) Locality is what makes a marginal-based decomposition possible at all.

CredalJT introduces, for each cluster C of a clique tree T , a vector $q_C \in \Delta(2^C)$ and optimizes the *cluster-feasible region*

$$\mathcal{G} = \left\{ (q_C)_{C \in T} : \sum_i q_C[i] = 1; \ M_{C \rightarrow S} q_C = M_{D \rightarrow S} q_D \text{ on every edge } (C, D), S = C \cap D; \text{ each hosted constraint } \right\} \quad (3)$$

The separator equalities $M_{C \rightarrow S} q_C = M_{D \rightarrow S} q_D$ force the two clusters to agree on the marginal of the shared atoms S ; a constraint is *hosted* on a cluster $C_h \supseteq$ scope and imposed on q_{C_h} 's marginal there. The read-off of $P(a \mid e)$ uses one cluster containing a and the evidence atoms.

5.6.2 Applicability: any chain graph, at the cost of its treewidth

The clique tree is built by the standard *moralize* $\rightarrow$ *triangulate* $\rightarrow$ *clique-tree* pipeline (`_node_scopes` + `min_fill_order` + `_jt_from_scopes`), which the package already uses for CCTE. The interaction graph is the moral graph of the chain-graph factorization (each family $\{\text{child}\} \cup \text{parents}$ becomes a clique, marrying co-parents); `min_fill_order` triangulates it; `_jt_from_scopes` returns the bucket tree, which is a clique tree with the running-intersection property (RIP). This pipeline is *topology-agnostic*: it runs on a chain, a polytree, a collider, a multiply-connected DAG, or a chain graph with undirected cliques alike. What varies is only the size of the largest cluster, i.e. the induced treewidth w ; the NLP has $\sum_C 2^{|C|} = O(n 2^{w+1})$ variables.

CredalJT therefore *applies to any chain-graph LCN*. The chain (Section 5.5.2) is merely the case $w=1$, where the clique tree is a linear path. On non-chain structures the clique tree branches: Table 3 shows that the engine builds a branching width-1 *tree* on **tree** (a single-parent branching process), a width-2 tree on **cancer** (collider $D \mid B \wedge C$) and on **polytree** (two colliders), a width-2 tree on **asia** (colliders + XOR formulas) and on **alarm** (compound $C - D$), and a width-5 *star* on **smokers** (the undirected clique $F_1 - F_2 - F_3$). The only role the budget $k_{\max}$ plays is to fall back to **ExactInference** when w is so large that $2^w \approx 2^n$ anyway (no saving to be had). The tree and the polytree are worked in detail in Section 5.6.5.

5.6.3 Exactness theorem

Let G^* be the chordal graph whose maximal cliques are the clusters of T (two atoms are adjacent iff they co-occur in some cluster). Say an LMC assertion $(X \perp Y \mid S)$ is *separator-entailed* if S separates X from Y in G^* .

Theorem 1 (Exactness). *Suppose*

- (H1) *T has the running-intersection property* (always true for the clique tree built above);
- (H2) *every LCN sentence is hosted: some cluster contains $\text{atoms}(s)$* ;
- (H3) *every LMC assertion is either hosted or separator-entailed*;
- (H4) *the query atom a and all evidence atoms lie together in some cluster*.

*Then for every such a , $\min_{q \in \mathcal{G}}(\text{read } a) = \min_{p \in \mathcal{F}} p_a$ and likewise for $\max$ (and for the conditional ratio $P(a \mid e)$). That is, **CredalJT**'s bounds equal the exact **ExactInference** bounds.*

Network	chain-graph feature	clique tree (atom clusters)	max cluster	w	exact?
chain	Markov chain ($n=8$)	linear path, $7 \times \{x_i, x_{i+1}\}$	2	1	✓
tree	branching tree ($n=7$)	6 edge-clusters $\{x_p, x_c\}$	2	1	✓
polytree	colliders ($n=7$)	path $\{x_0x_1x_2\}-\{x_2x_3\}-\{x_3x_4x_5\}-\{x_5x_6\}$	3	2	✓ [†]
d4_biting	cross-family sentence	single cluster $\{A, B, C\}$	3	2	✓
cancer	collider $D \mid B \wedge C$	tree, e.g. $\{A, B, C\}-\{B, C, D\}-\{D, E\}$	3	2	✓
asia	colliders + XOR	tree, $\{B, S, C\}-\{B, C, D\}-\{C, D, X\}$	3	2	✓ [†]
alarm	compound node $C-D$	tree, $\{B, E, A\}-\{A\}-\{A, C, D\}$	3	2	✓
smokers	undirected clique $F_1F_2F_3$	star, centre separator $\{S_1, S_2, S_3\}$	6	5	<i>outer only</i>

Table 3: **CredalJT** builds a genuine (branching) junction tree on every chain-graph LCN, not only on chains. “Max cluster” is the largest cluster’s atom count (what the engine reports as `induced_width`); the treewidth w is one less, and governs cost ($O(n2^{w+1})$ variables). The last column is the exactness verdict of Section 5.6.3: ✓ where the sufficient condition holds; *outer only* on **smokers**, where six large-scope LMC assertions are neither hosted nor separator-entailed. [†]**asia** and **polytree** *violate* the sufficient condition (unentailed, unhostable LMC drops — the co-parent independencies of their colliders) yet are exact on every queried atom: the condition is sufficient, not necessary (Sections 5.6.3, 5.6.5).

Proof. We show the two feasible regions have identical projections onto every cluster scope; since each objective (a linear functional of one cluster’s marginal, or a ratio of two such) depends only on that projection, the optima coincide.

Projection $\mathcal{F} \rightarrow \mathcal{G}$ (*always valid, hence an outer bound*). Given $p \in \mathcal{F}$, set $q_C := p \upharpoonright C$ for each cluster. Cluster simplexes hold (a marginal of a distribution sums to 1); separator equalities hold (both sides equal $p \upharpoonright S$); every hosted constraint c holds because $q_{C_h} \upharpoonright \text{scope}(c) = p \upharpoonright \text{scope}(c)$ and c holds for p by *locality*. The read-off of a equals p_a . Hence $q \in \mathcal{G}$ with the same objective value, so $[\min, \max]_{\mathcal{G}} \supseteq [\min, \max]_{\mathcal{F}}$: **CredalJT is always a valid outer bound**, with no assumption beyond (H1)/(H4).

Reconstruction $\mathcal{G} \rightarrow \mathcal{F}$ (*needs H2/H3*). Given a separator-consistent family $(q_C) \in \mathcal{G}$, the RIP marginal-consistency theorem (Lauritzen, *Graphical Models*, Prop. 3.18) yields a *unique* joint

$$p = \frac{\prod_C q_C}{\prod_{(C,D) \in T} q_S} \quad (4)$$

that is a genuine distribution and whose marginal on each cluster C is exactly q_C (and hence on each separator and on each sub-clique). We verify $p \in \mathcal{F}$:

- *Sentences* (H2): $p \upharpoonright \text{atoms}(s) = q_{C_h} \upharpoonright \text{atoms}(s)$ for the host C_h , which satisfies g_s by construction of $\mathcal{G}$; by locality $g_s(p)$ is in range.
- *Hosted LMC* (part of H3): identical argument — the bilinear residual is a function of the host cluster’s marginal, which satisfies it.
- *Separator-entailed LMC* (the other part of H3): p in (4) factorizes over T , so it obeys the *global Markov property of G^** : whenever S separates X from Y in G^* , $X \perp_p Y \mid S$. Thus every separator-entailed LMC holds for p *automatically*, even though it was dropped from $\mathcal{G}$.

So $p \in \mathcal{F}$ with p_a equal to the read-off, giving $[\min, \max]_{\mathcal{F}} \supseteq [\min, \max]_{\mathcal{G}}$.

Both inclusions force equality. For a conditional query the objective is the ratio of two marginal functionals of the single cluster of (H4); the same two maps preserve it. $\square$ $\square$

Reading of the proof. (a) The argument never mentions chains: it uses only RIP, hosting, separator-entailment, and locality. Exactness is a property of the *junction tree*, not of the chain. (b) The projection half holds unconditionally, so **CredalJT** is *never unsound* — its interval always contains the true one; (H2)/(H3) only decide whether it is *tight*. (c) On a pure Markov chain the moral graph is already triangulated and every LMC $(x_{i+1} \perp x_{<i} \mid x_i)$ is a single separator, so (H3) holds with $w=1$ — recovering Section 5.5.2 as the special case.

What the engine actually checks, and the residual heuristic. **CredalJT** hosts every sentence (H2) and every *hostable* LMC, and it recognises a *single-edge* separator-entailed LMC via `_rip_implied` and drops it. The gap between this test and the full (H3) condition is that `_rip_implied` only inspects *one* tree edge, whereas separator-entailment may rest on the chordal graph G^* as a whole (multiple cliques). For an LMC that is neither hostable nor single-edge-RIP, the engine *drops it anyway* — this is the “inert drop”. The drop is exactness-preserving *iff* the assertion is separator-entailed by G^* ; when it is not, the result is a valid *outer* bound that can be loose. (The earlier description of these drops as “inert because a product joint satisfies every LMC” is the right intuition for the *strong-extension* product but is not the reason exactness holds: $\mathcal{F}$ is *not* the product set; the correct reason is global-Markov entailment by G^* .)

5.6.4 Empirical validation across structures

We checked all six shipped chain-graph LCNs two ways: a solver-free structural test (does each LMC satisfy (H3)? — i.e. separation in G^* , computed from the clusters), and a numerical comparison of **CredalJT** against an *independent* exact oracle. For the ≤ 5 -atom networks the oracle is the LMC-aware brute-force verifier (`verify_bruteforce.py`), which optimizes over a parametric family that satisfies every LMC equality *by construction* (residual $\sim 10^{-16}$) using SLSQP only — so it is immune to the nonconvex-solver unreliability that makes raw **ExactInference**/SCIP themselves untrustworthy on these bilinear systems. The results:

- **chain, d4_biting, cancer:** (H2)/(H3) hold (every LMC is hosted or single-edge-RIP), so Theorem 1 guarantees exactness; the numbers match the oracle on every atom.
- **alarm:** the two five-atom LMCs $(C \perp B, E \mid A, D)$, $(D \perp B, E \mid C, A)$ are not single-edge-RIP but *are* separator-entailed by G^* ($\{A, D\}$ separates C from $\{B, E\}$), so (H3) holds and the theorem applies. The brute-force oracle confirms the exact bounds $P(A)=[0.087, 0.100]$, $P(B)=[0.100, 0.123]$, $P(E)=[0.050, 0.074]$, $P(C)=P(D)=[0, 1]$; **CredalJT** reproduces them, whereas raw **ExactInference** (local solver) reports the spuriously-tight $P(A)=[0.092, 0.100]$, $P(E)=[0.050, 0.073]$ — **CredalJT** is here the *more* trustworthy engine.
- **asia:** three LMCs $((B \perp C, X \mid S), (D \perp S, X \mid C, B), (X \perp S, B, D \mid C))$ are *neither* hosted *nor* separator-entailed, so (H3) *fails* and the theorem does not certify exactness. Empirically **CredalJT** is nonetheless exact on every queried atom: the untailed assertions happen not to bind those marginals given the sentences. This shows (H3) is *sufficient but not necessary*.
- **smokers:** all six large-scope LMCs (e.g. $(S_3 \perp F_1, C_1, C_2 \mid S_1, S_2, F_2, F_3)$) are unhostable and untailed — in the width-6 family cluster $\{F_1, F_2, F_3, S_1, S_2, S_3\}$ the atoms S_3 and F_1 co-occur, so no separator can divide them. (H3) fails *and* it bites: **CredalJT** returns the loose $P(F_1)=P(S_1)=[0, 1]$ where the dropped independencies would tighten them. The interval is still a valid *outer* bound (the projection half of the proof), just not tight.

The practical upshot: **CredalJT** is a *general* chain-graph engine whose bounds are **always sound (outer)** and **exact whenever (H2)/(H3) hold** — which provably covers a clean chain, a branching *tree* (single-parent nodes, treewidth 1; Section 5.6.5), **cancer**, and **alarm**. It is *also* exact empirically on the structures that violate (H3) but whose dropped independencies do not bind — **asia** and the *polytree* of Section 5.6.5. The one failure mode is an LMC whose scope is buried inside a single large cluster (an undirected clique, as in **smokers**), where no separator can express it and it is too large to host; there the engine silently widens to an outer bound. A conservative variant would *refuse to drop* any non-separator-entailed LMC and instead widen the host (or fall back to **ExactInference**) so that the output is always exact; the current engine trades that guarantee for treewidth, and the `d5_exact` flag together with this G^* -separation test tells the caller which regime it is in.

5.6.5 Worked structures: a branching tree and a polytree

The chain of Section 5.5.2 is the simplest singly-connected structure. Two natural generalizations stress the exactness theorem in different ways: a *tree* (each node still has a single parent, but the graph *branches*) and a *polytree* (the undirected skeleton is still a tree, but nodes may have *several* parents, i.e. COLLIDER v-structures). We ship both as `examples/tree.lcn` and `examples/polytree.lcn` and analyse them here. The headline: **the tree is exact for the same reason the chain is** (Theorem 1 applies verbatim, treewidth 1); **the polytree is exact too, but only empirically** — moralizing its colliders breaks the sufficient condition (H3), yet the dropped independencies turn out not to bind, exactly as for **asia**.

The branching tree (tree.lcn). Seven atoms in a binary tree

$$x_0 \rightarrow \{x_1, x_2\}, \quad x_1 \rightarrow \{x_3, x_4\}, \quad x_2 \rightarrow \{x_5, x_6\},$$

a prior on the root x_0 and a transition $P(x_c | x_p)$ on each edge (no cross-family sentence). Because *every node has exactly one parent*, there are no co-parents to marry: the moral graph *equals* the tree skeleton, already triangulated. **CredalJT** builds one two-atom cluster per edge,

$$\{x_0, x_1\}, \{x_0, x_2\}, \{x_1, x_3\}, \{x_1, x_4\}, \{x_2, x_5\}, \{x_2, x_6\}$$

(six clusters, single-atom separators $\{x_0\}, \{x_1\}, \{x_2\}$, plus a trivial $\{x_6\}$ leaf), giving **treewidth 1**, just like the chain but with a branch. Every LMC assertion ($x_c \perp \text{non-descendants} | x_p$) is then a single-separator (RIP) independence — $\{x_p\}$ separates the child’s subtree from the rest in G^* — so all are dropped as RIP-implied and *none* is imposed, while (H2)/(H3) hold trivially. Theorem 1 therefore certifies exactness, and the brute-force oracle confirms it: **CredalJT** and the LMC-correct oracle agree to four decimals on all seven atoms, e.g.

$$P(x_0) = [0.300, 0.500], \quad P(x_1) = P(x_2) = [0.400, 0.680], \quad P(x_3) = \dots = P(x_6) = [0.328, 0.640],$$

solving an NLP with $\sum_C 2^{|C|} = 6 \cdot 2^2 = 24$ variables rather than the full joint $2^7 = 128$. *A branching tree is exactly as favourable as a chain.*

The polytree (polytree.lcn). Seven atoms, singly connected, with two colliders:

$$x_0 \rightarrow x_2 \leftarrow x_1, \quad x_2 \rightarrow x_3, \quad x_3 \rightarrow x_5 \leftarrow x_4, \quad x_5 \rightarrow x_6,$$

priors on the roots x_0, x_1, x_4 and the four-row conditionals $P(x_2 | x_0, x_1), P(x_5 | x_3, x_4)$. The multi-parent nodes are what distinguishes a polytree from a tree, and they change the Markov picture

in two ways. First, the LMC now contains *marginal*-independence assertions of the co-parents and unrelated roots — $(x_0 \perp x_1, x_4)$, $(x_1 \perp x_0, x_4)$, $(x_4 \perp x_0, x_1, x_2, x_3)$ — and a small conditional one $(x_2 \perp x_4 \mid x_0, x_1)$. Second, moralizing each collider *marries its co-parents*: the family cliques become $\{x_0, x_1, x_2\}$ and $\{x_3, x_4, x_5\}$, so **CredalJT** runs at **treewidth 2**, with the clique tree

$$\{x_0, x_1, x_2\} \overset{\{x_2\}}{-} \{x_2, x_3\} \overset{\{x_3\}}{-} \{x_3, x_4, x_5\} \overset{\{x_5\}}{-} \{x_5, x_6\}.$$

(7 clusters, max cluster 3 atoms; $\sum_C 2^{|C|} = 2 \cdot 2^3 + \dots = 28$ variables vs. $2^7 = 128$.)

Now the subtlety. The co-parent assertion $(x_0 \perp x_1, x_4)$ has scope $\{x_0, x_1, x_4\}$, which straddles *both* collider cliques and so fits in *no* single cluster — it is **unhostable**. And after moralization x_0 and x_1 co-occur in $\{x_0, x_1, x_2\}$, so no separator can divide them: it is also **not** separator-entailed by G^* . The same holds for all five co-parent / root independencies. **CredalJT** therefore *drops* them (imposes none of the LMC assertions), and **condition (H3) of Theorem 1 fails** — the theorem does *not* certify the polytree exact. This is structurally the **smokers** situation (an unentailed, unhostable independence buried in a moralized clique), now arising on a clean polytree rather than from an undirected clique.

Yet the brute-force oracle (a parametric family $P(x_0)P(x_1)P(x_4)P(x_2 \mid x_0, x_1)P(x_3 \mid x_2)P(x_5 \mid x_3, x_4)P(x_6 \mid x_5)$, LMC-correct to 3×10^{-16}) shows **CredalJT** is in fact **exact on every atom**:

	x_0	x_1	x_2	x_3	x_4	x_5	x_6
CredalJT (D5, SCIP)	[.30, .50]	[.40, .60]	[.31, .63]	[.348, .676]	[.20, .50]	[.264, .653]	[.339, .694]
brute-force oracle	[.30, .50]	[.40, .60]	[.31, .63]	[.348, .676]	[.20, .50]	[.264, .653]	[.339, .694]

The dropped co-parent independencies simply do not bind these marginals: although the joint cluster $\{x_0, x_1, x_2\}$ is *permitted* to correlate x_0 and x_1 , the optima of every singleton marginal $P(x_i)$ over the cluster NLP are attained without exploiting that freedom (the optimizing joint happens to keep the co-parents independent). This is the polytree analogue of the **asia** observation: **(H3) is sufficient, not necessary**, so a structure can violate it and still be exact. (One can manufacture a polytree where the drop *does* bite — add a cross-family sentence coupling two co-parents, e.g. $P(x_0 \wedge x_1) \leq \epsilon$, whose scope sits in the collider clique and would then be hosted, or one coupling a co-parent to a different collider, which would be unhostable and make the bound loose — but the plain transition-only polytree above is exact.)

Takeaway. A tree is exact and cheap (treewidth 1) by the same RIP argument as the chain. A polytree is more delicate: its colliders moralize into cliques that hide the co-parent independencies from the separators, so the engine drops them and the *sufficient* exactness condition no longer holds — the bound is then guaranteed only to be a valid *outer* bound. On the transition-only polytree we ship it remains exact because those independencies do not bind, but the certificate is empirical (oracle agreement), not the theorem. The reliable rule of thumb is the G^* -separation test of Section 5.6.3: a singly-connected structure with *single-parent* nodes (chain, tree) always satisfies (H3) and is provably exact at treewidth 1; a structure with *multi-parent* colliders (polytree) may need the empirical check, and a buried undirected clique (**smokers**) is where the drop actually loosens the bound.

The **cn/** package ships five inference algorithms beyond CVE. Crucially, *all of them consume the same object*: a built **CredalNetworkVertices**, i.e. the directed structure plus the enumerated extreme points of the local credal sets (**cnv.extreme_points**). They differ only in how they propagate those vertices:

Algorithm	Propagation	Bound character vs. the strong extension
Credal VE (<code>cve.py</code>)	per-target bucket elimination	exact (the strong-extension interval)
Credal CTE (<code>ccte.py</code>)	two-pass bucket tree	exact ($\epsilon=0$); outer if $\epsilon>0$
ApproxLP (<code>approxlp.py</code>)	per-variable LP over vertices	<i>inner</i> (subset of the true interval)
Interval BP (<code>ibp.py</code>)	loopy interval messages	<i>outer</i> (interval); inner (variational)
Credal IJGP (<code>ijgp.py</code>)	join-graph, <i>i</i> -bound	outer-style; tunable by <i>i</i> -bound

(`CredalVE.run(evidence)` now computes *all* singleton-atom marginals in one call by looping the single-query bucket elimination over the credal-network nodes: for each target node it recomputes the elimination order so the target is eliminated last, runs one pass, and reads off the node’s bounds, projecting compound nodes to their atoms by an LP. Because each target gets its own optimal per-target order, its strong-extension bound is at least as tight as Credal CTE’s shared two-pass schedule.)

The essential observation is that every one of these algorithms brackets the *strong-extension* interval $[q_{SE}, \bar{q}_{SE}]$, not the LCN interval $[q_{LCN}, \bar{q}_{LCN}]$. Since

$$[q_{LCN}, \bar{q}_{LCN}] \subseteq [q_{SE}, \bar{q}_{SE}],$$

shrinking the strong extension toward $\mathcal{D}_{LCN}$ pulls every algorithm’s output inward. This splits the schemes into two leverage classes.

Class 1 — input-level tightening (D1, D2, D3): free for all five algorithms. D1/D2/D3 change only what `CredalNetworkVertices.from_lcn` produces—tighter local intervals, hence tighter enumerated vertices. Because the algorithms read those vertices without assuming how they were obtained, *no algorithm code changes*. The effects compose with each algorithm’s own character:

- **CVE / CCTE** (exact over the strong extension): their interval is exactly $[q_{SE}, \bar{q}_{SE}]$, so a tighter strong extension is reported verbatim—the cleanest beneficiary.
- **ApproxLP** (inner): its interval sits *inside* the strong-extension interval. Tighter vertices raise its lower bound and lower its upper bound too, but the inner gap remains; pairing ApproxLP (inner) with CVE (outer) on the *same* tightened vertices brackets the LCN truth from both sides.
- **Interval BP** (outer): tighter local messages mean tighter fixed-point intervals; the loopy outer slack is unchanged but starts from a narrower base.
- **IJGP**: tighter cluster potentials propagate through the join graph; the *i*-bound still trades cost for tightness independently.

On the example, replacing the vacuous $[0, 1]$ rows (D1) with genuine sub-intervals narrows the vertex sets fed to *all five*, so e.g. a query $P(A)$ tightens for CVE, ApproxLP, IBP, CTE and IJGP simultaneously, with one build-stage change.

Class 2 — propagation-level coupling (D4, D5): per-algorithm work. D4/D5 do not merely narrow intervals; they remove vertex *combinations* (D4) or avoid forming the product altogether (D5). This is a change to the propagation, so each algorithm needs its own adaptation:

- **CVE / CCTE**: add the cross-vertex side constraints (D4) to the bucket / cluster combination step. This is the most natural fit—both already optimize over vertex tuples per bucket, so a coupling constraint slots in directly. D5 is essentially a constraint-aware rewrite of CVE’s bucket optimization.
- **ApproxLP**: its per-variable LP already optimizes over a vertex polytope; D4’s coupling constraints can be added as extra LP rows, keeping the inner-bound guarantee while shrinking the polytope.
- **Interval BP / IJGP**: message passing does not carry global cross-factor constraints naturally. The practical route is to let D4’s coupling be absorbed at the *build* stage where possible (so it reduces to Class 1 for these two), and accept that the residual coupling D4 cannot localize is simply not captured by loopy / join-graph propagation—these remain the loosest of the family, by design.

Practical recommendation. Because Class 1 is free for all algorithms, apply D1 (and D2 where affordable) at the build stage *first*; every downstream algorithm benefits without modification. Reserve the per-algorithm D4/D5 work for CVE/CCTE (and ApproxLP), which can absorb the coupling exactly, and treat IBP/IJGP as fast outer approximations that ride on the Class-1 tightening.

6 Validation strategy

Tightening must be checked against ground truth on two fronts, matching the two gaps.

Local-set tightness (gap A) via D3. D3 is itself the certifier: for any candidate local credal set produced by D1 or D2, compare its interval $[p, \bar{p}]$ for each (family, interpretation) against the D3 certified-global interval $[p^*, \bar{p}^*]$ obtained from the SCIP joint-LP. A correct outer approximation requires $[p^*, \bar{p}^*] \subseteq [p, \bar{p}]$, and the width difference quantifies the residual local looseness D1/D2 leave behind. Because D3 uses the global `solver="scip"` path, these reference intervals are certified, not merely local optima—unlike `ExactInference`, which is only a local solver and is not a reliable oracle.

Extension tightness (gap B) via `verify_lmc_factorization.py`. The factorization verifier reconstructs each product joint from a choice of local extreme points and evaluates the LMC residual

$$P(x, y, s) P(s) - P(x, s) P(y, s) \approx 0$$

for every assertion (its `lmc_constraint_groups_vec` residuals). A residual that stays $\leq \text{tol}$ after a tightening change confirms the scheme has not broken the independence structure while shrinking the set; a non-trivial residual that the LCN does *not* permit localizes a strong-extension (gap B) violation that D4/D5 must address. This check is solver-free and deterministic, so it is a stable regression gate.

End-to-end. On small instances, posterior marginals from CVE under each scheme are compared to the D3-style full-joint bounds for the query atom (the same certified global solve, with the query as objective). The recommended sequence is D1, then D2 with D3 as the per-instance certifier of gap A; D4/D5 are pursued only once the D3 comparison shows the residual gap is genuinely the strong-extension gap B rather than remaining local looseness.